

Measuring the Value of AI in Your Bank

How to measure staff capacity, operating cost, and customer value without overstating the return

Useful capacity. Credible evidence. Better decisions.

An AI business case becomes stronger when it explains what changed, who benefited, and what the bank actually paid. This guide shows how to measure the complete workflow, distinguish capacity from cash savings, and decide whether the result deserves further investment.

01

Measure the complete job

02

Separate capacity from cash

03

Build an evidence-based decision

For bank CFOs, COOs, finance teams, and business owners

An educational guide with a worked banking example and a short supporting worksheet.

01 / DEFINE VALUE

A faster draft is not the whole result.

AI can create useful capacity without reducing payroll. It can also produce more output while making the complete process slower. A credible business case distinguishes those outcomes.

Begin with a business question: which recurring job should become better, for whom, and under what quality standard? A demonstration can show that a draft is possible. It does not establish that employees can use it, that review is affordable, or that a customer receives a better result.

Value category	What would support the claim	What it does not mean
Staff capacity	Less total staff effort for comparable, acceptable completed work.	An automatic payroll reduction or a reduction in headcount.
Cash impact	A verified expense reduction or incremental contribution, compared with incremental cash costs.	The dollar value of every hour an employee could theoretically reuse.
Service outcome	Evidence of shorter queues, timely follow-up, or other useful customer or staff outcomes.	A faster AI response or a higher count of generated documents.
Quality and control	Fewer substantiated defects, dependable review, and acceptable information handling.	A guarantee, regulatory approval, or a dollar value for hypothetical losses avoided.

Choose the primary benefit before the pilot

A service-led pilot might aim to release operations staff for overdue branch support. A cost-led pilot might target a documented overtime expense. Either can be legitimate, but the evidence and financial treatment differ. Do not switch the definition after seeing which metric looks most impressive.

Use a chain of evidence.

Comparable work → acceptable quality → lower total effort → usable staff capacity → observed service or cash outcome. A missing link remains an assumption, not a result.

The interagency community-bank guide encourages evaluating strategic fit, direct and indirect costs, expected benefits, and risks. NIST’s voluntary AI framework calls for evaluation in the intended context. Neither source supplies a universal AI ROI formula. The calculations in this guide are an educational method, not an agency requirement. [1, 2]

Reading path: measurement foundations, pages 2–6; fictional worked example, pages 7–10; sensitivity and decisions, pages 11–12; supporting worksheet, page 13; references, page 14. All example figures are invented teaching assumptions, not observed bank or product results.

02 / BUILD A FAIR BASELINE

Measure one complete unit of work.

A comparison is useful only when the old and proposed processes are solving substantially the same problem to the same standard.

Define the unit before starting the clock. For an internal policy-support workflow, it might be one staff question answered with an approved source, checked by a qualified reviewer, and returned in usable form. It is not one prompt, one search, or one AI response. Several attempts can belong to the same completed request.

Set this boundary	Why it matters
Start and finish	Begin when a suitable request is ready for handling. End when the agreed, reviewed deliverable is complete. Record later corrections separately and add their effort.
Eligible work	Define source access, task type, permitted information, required approval, and exclusions. Keep out-of-scope cases visible rather than silently removing difficult work.
Complexity	Group routine questions, conflicting-source questions, and exceptions separately. A different mix can explain an apparent productivity gain.
Volume and period	Use the same unit and reporting period for both processes. Distinguish actual measured counts from forecasts or modeled monthly volumes.
People and conditions	Record roles, experience, source quality, interruptions, service demand, and relevant system changes. Do not compare experts with new users without explaining the difference.

Compare rates without losing the workload

First compare staff minutes per completed request within each task category. Then apply those rates to a common workload mix. If routine questions dominate the pilot but difficult questions dominate the baseline, the unadjusted averages are not comparable. Also report total hours: lower effort per request can coexist with higher total effort when demand increases.

Include the current process as it actually operates, including ordinary review and correction. Do not use an artificially slow manual process as the baseline. Compare with a simpler alternative, such as better templates, search, or a clarified procedure, if it could solve the same problem at lower cost.

Keep the denominator honest.
Report all eligible requests, the subset on which AI was attempted, the subset completed through the assisted path, and the subset sent back to manual handling. Adoption and successful completion are different rates.

A sample must reflect the conditions for which a claim will be made. NIST emphasizes deployment-relevant evaluation, documented limitations, and performance benchmarks. A small or unusually easy sample supports only a limited conclusion. [1]

03 / COUNT THE WHOLE CLOCK

Add every person’s effort, not the wait.

Staff time, elapsed turnaround, and unattended machine time answer different questions. Track them separately before turning minutes into a benefit.

Time measure	What to include	How to use it
Active staff effort	Preparing sources, operating tools, checking answers, correcting work, escalation, and required active supervision across all roles.	Use it to estimate staff capacity. Two people working together for 10 minutes consume 20 staff minutes.
Elapsed turnaround	Time from request readiness to accepted completion, including queues, handoffs, and waits.	Use it to evaluate service. Do not multiply every elapsed minute by a wage rate.
Unattended processing	Machine work that does not require an employee’s attention or prevent other useful work.	Track reliability, throughput, and service cost. It is not automatically paid staff effort.
Recurring overhead	Source maintenance, routine coordination, usage review, support, and continuing monitoring.	Deduct staff effort once per reporting period unless already included in task timings.

Write the time rule before collecting data

Use task logs, short activity timers, or structured observations rather than relying only on recollection. Record genuine waiting constraints: if someone must monitor a process and cannot do other work, that attention belongs in staff effort. Do not count a background run and a simultaneous productive task as two full blocks of the same employee’s time.

Net recurring capacity = baseline staff hours – assisted-process staff hours – separate recurring administration hours.

The assisted-process total includes completed assisted work, failed attempts followed by manual work, and eligible work that never used AI. One-time setup and initial training belong in the investment record, not hidden inside a steady-state monthly rate. A ramp-up period needs its own measured effort and volume.

Preserve the role-level view

Ten minutes removed from operations can become ten minutes of specialist review. Hours are unchanged, but cost and bottlenecks may differ. Keep hours by role and use Finance-approved rates. The worked example uses a blended rate only to simplify the arithmetic.

Count each cost once, in the right view.

If administration already reduces net hours, do not subtract its labor value again from a capacity calculation. A separate cash expense for an outside monitoring service is different work and may still belong in cash costs. Label both explicitly.

04 / MAKE QUALITY A CONDITION

Faster is useful only when the work holds up.

The time comparison should not reward an answer that skips evidence, shifts corrections downstream, or creates an unacceptable information-handling failure.

Ask a domain owner to define an acceptable deliverable before testing. For policy support, that can include the right source, its approval status, applicable wording, a faithful interpretation, and a clear statement of unresolved questions. A well-written answer that cites the wrong procedure does not pass because it was quick.

Measure	Define the numerator and denominator	What to retain
First-review acceptance	Drafts accepted without substantive correction ÷ drafts reviewed.	Criteria, counts, reviewer notes, and excluded cases.
Correction burden	Total correction minutes and correction types across all attempted requests.	Who corrected the work and whether another request was delayed.
Fallback rate	Attempted requests abandoned for manual completion ÷ all AI-attempted requests.	Failed-attempt effort, full manual effort, and reason.
Evidence accuracy	Checked citations or extracted items that meet the agreed standard ÷ all checked items.	Source version, the checked passage or field, and severity of each defect.
Control exceptions	Count and describe unacceptable access, disclosure, or unapproved actions separately.	Incident facts and required response, not merely an average score.

Do not let the average conceal the consequence

Report both the frequency and severity of defects. A formatting correction and a materially misleading policy interpretation should not disappear into one error count. Keep quality thresholds tied to the use and its consequences. An unacceptable disclosure or action triggers the bank’s established response even if the aggregate time result is favorable.

Check a representative share of accepted work and investigate reported problems after handoff. If you only inspect rejected drafts, you cannot tell whether reviewers missed problems in accepted ones. Protect test records and source material under the bank’s approved access and retention rules.

A clean sample is not a guarantee.

Report the sample size and tested conditions. “No critical failures observed in this sample” is a bounded statement. “The system makes no critical errors” is not supported by that observation.

The criteria here are proposed measurement practices, not regulatory thresholds. NIST calls for documented evaluation, uncertainty and limitations, and monitoring of performance and emerging risks. [1]

05 / RUN A CREDIBLE COMPARISON

Design the pilot before reading the results.

A simple written protocol prevents a persuasive demonstration from becoming the entire business case.

1. Name the claim and its owner

State the intended workflow, business owner, quality criteria, information boundaries, reporting period, and decision to be made. Separate a service objective from a cash objective. Technology and risk colleagues determine the required operating reviews; Finance owns the treatment of costs and benefits.

2. Choose a comparison that can be explained

Where feasible, compare equivalent task groups handled under both processes during similar conditions. Random assignment can reduce selection bias when it is practical and appropriate. If using the same synthetic cases twice, counterbalance the order or use independent handlers so the second attempt does not benefit mainly from remembering the answer. A before-and-after comparison is weaker when demand, staffing, or procedures also change.

3. Separate learning from evaluation

Record initial training, configuration, and experimentation as setup. Keep evaluation examples separate from cases used to tune instructions. Do not remove early failures without reporting them; explain when a stable operating configuration was reached and why the reporting window was chosen.

4. Collect enough detail to challenge the conclusion

For each request, record a non-sensitive ID, task category, approved source version, assigned path, AI attempts, staff minutes by role, review outcome, corrections, fallback, and elapsed completion time. Summarize counts, means, medians, and the spread of results. Look separately at difficult cases and unusually slow requests.

5. Review the evidence, including what did not improve

There is no universal sample size that proves an AI workflow is valuable. Choose coverage around variability, consequence, and the decision at stake, with analytical help for material investment claims. Document uncertainty instead of manufacturing precision. Recheck results after meaningful changes to the sources, configuration, staffing, or scope. [1]

Keep observation, estimate, and target separate.

Observed: logged results from the evaluated work. Estimated: an extrapolation using stated assumptions. Target: a desired future result. Put those labels in the decision memo so a forecast cannot be mistaken for a delivered benefit.

The numerical example on the next pages is a hypothetical monthly model. It does not represent a completed pilot, a measured sample, or an expected result from any specific product.

06 / FICTIONAL EXAMPLE: STAFF CAPACITY

Follow all 300 requests through the process.

A fictional community bank is evaluating internal policy support using approved procedures, no customer records, and human review. Every number below is a teaching assumption.

The common workload is 300 eligible requests per month. The existing process takes 30 staff minutes per request across all people involved. The proposed process attempts AI on 240 requests, or 80%. It completes 216 through the assisted path and sends 24 back to the full manual process. Another 60 never use AI.

Proposed-process path	Requests	Staff minutes each	Monthly hours
Reviewed assisted completion	216	16 = 6 preparation + 8 review + 2 correction	57.6
Abandoned attempt, then manual	24	38 = 8 attempted + 30 manual	15.2
No AI attempt; ordinary manual work	60	30	30.0
Separate recurring administration	All work	6 hours per month	6.0
Total proposed staff effort	300	Includes every path	108.8

Baseline: $300 \times 30 \div 60 = 150.0$ staff hours.

Net recurring capacity: $150.0 - 108.8 = 41.2$ hours per month.

The denominator changes the story

AI was attempted on 80% of eligible requests, but only 72% of all eligible requests finished through the assisted path: $216 \div 300$. The fallback rate is 10% of attempted requests: $24 \div 240$. These percentages describe different stages; none should be substituted for the others.

The average staff time across all 240 attempted requests is 18.2 minutes, not 16. It includes both assisted completions and fallbacks. Before administration, those attempts release 47.2 hours compared with manual handling. The separate 6 administration hours reduce that benefit to 41.2.

Why the optimistic shortcut fails

Using $300 \times (30 - 16) \div 60$ would claim 70 hours. It incorrectly assumes full adoption, no fallbacks, and no administration. Following the actual modeled paths gives 41.2 hours, before asking whether those hours can be used.

The 2 correction minutes are an average allowance on the completed assisted path, not a claim that every request needs identical correction. The model assumes comparable accepted quality; a real bank must test that condition. These are invented figures, not customer results.

07 / FICTIONAL EXAMPLE: COMPLETE COSTS

The subscription is not the full investment.

Separate cash outlays, staff opportunity cost, and the timing of each. That makes the calculation reviewable without pretending they are interchangeable.

Illustrative cost item	Amount	Treatment
Initial deployment and configuration	\$2,500	One-time cash
Additional hardware	\$1,200	One-time cash
External initial training	\$800	One-time cash
Total initial cash outlay	\$4,500	Paid at inception in the example
Licenses and service usage	\$500/month	Recurring external cash
Incremental infrastructure	\$150/month	Recurring external cash
Outside support	\$200/month	Recurring external cash
Independent monitoring service	\$100/month	Recurring external cash
Total recurring external cost	\$950/month	Excludes internal staff time

Show internal effort even when payroll is unchanged

Assume initial staff setup and training require 60 hours within existing paid schedules. At an illustrative \$52 blended loaded hourly rate, that is \$3,120 of opportunity cost, not \$3,120 of new cash spending. The recurring 6 administration hours are already included in the capacity calculation. Do not deduct them again there.

Year-one incremental cash costs: $\$4,500 + (12 \times \$950) = \$15,900$.

Separate initial staff opportunity cost: $60 \times \$52 = \$3,120$.

The example assumes the full recurring run rate begins in month one. A real first-year budget should show monthly setup, training, adoption, usage, and benefit ramp. Hardware replacement, paid incident support, higher usage tiers, contract increases, and exit work may also matter. Include them when relevant rather than treating this teaching budget as a complete vendor quote.

Cash flow is not the income statement.

Finance determines expense recognition, capitalization, depreciation, and tax treatment. Do not add depreciation to a cash model that already includes the same hardware purchase. An allocation of existing overhead is not automatically a new cash cost.

These prices are invented, not Providensi pricing. The community-bank guide explicitly includes direct contractual costs and indirect changes to staffing, processes, and technology in planning considerations. [2]

08 / FICTIONAL EXAMPLE: USEFUL REDEPLOYMENT

Reclaimed time still needs a destination.

Forty-one hours spread across short gaps, different employees, and different shifts do not automatically create a usable block of capacity.

In the service-led version of the fictional example, the operations manager identifies approved follow-up work that staff can perform in the released time. Suppose 75% of the modeled 41.2 hours can be usefully redeployed. That yields 30.9 hours; the other 10.3 hours remain uncommitted, fragmented, or otherwise not demonstrated as useful.

Useful redeployment: $41.2 \times 75\% = 30.9$ hours per month.

Illustrative capacity value: $30.9 \times \$52 = \$1,606.80$ per month.

The \$1,606.80 is an opportunity-value proxy at an assumed labor rate. It is not cash saved, profit, or new revenue. Valuing all 41.2 hours would produce \$2,142.40, but that would include time not shown as usefully redeployed. The two values overlap and must never be added together.

Intended destination	Evidence to collect	Avoid this shortcut
Branch support	Additional eligible follow-ups completed, aging of the queue, and staff time used.	Calling every free minute improved service.
Procedure maintenance	Reviewed updates completed, remaining backlog, and an accountable content owner.	Counting drafts without checking acceptance.
Customer attention	Timely callbacks or other approved service work, with quality and complaint measures.	Attributing every later sale to AI.
Reduced overtime	Actual paid hours removed at an avoidable cash rate, with scheduling and service evidence.	Adding the same hours to both service and cash benefits.

An outcome claim needs more than an activity count

Track whether the useful work changes the problem the sponsor selected. More follow-ups are a process measure; a sustained reduction in overdue work is an outcome. Compare against demand, seasonality, staffing, and other changes. If an extra marketing campaign or improved procedure also contributed, explain that rather than assigning the whole difference to AI.

Be especially careful with revenue.

Loan principal, deposits, and gross sales are not incremental profit. A revenue case needs attributable incremental contribution after relevant costs, with Finance validation. Do not add that contribution to a labor-value proxy for the same activity without addressing the overlap.

In practice, verify role-level availability and rates instead of assuming one blended rate fits every employee. The example's 75% redeployment assumption is not an industry benchmark or a forecast.

09 / FICTIONAL EXAMPLE: CASH AND RETURN

A cash-savings claim needs cash evidence.

The service-led case can be worthwhile without producing a positive cash return. A separate cost-led case requires an expense that genuinely disappears.

Base case: capacity improves, but spending does not fall

In the base case, payroll and other existing expenses remain unchanged, and no new cash benefit has been demonstrated. Year-one incremental cash flow is therefore negative \$15,900. The bank may still value better service or capacity, but it should not describe the opportunity-value proxy as cash ROI.

Alternative case: documented overtime is actually removed

Instead of the service-redeployment use on page 9, assume the same capacity allows 28 overtime hours per month to be eliminated at a genuinely avoidable cash rate of \$50 per hour. Finance verifies the lower paid amount, appropriate skills and scheduling, unchanged service quality, and no replacement cost elsewhere. This is a separate hypothetical case, not an added benefit.

Simple cash measure	Calculation
Monthly cash benefit	$28 \times \$50 = \$1,400$
Monthly net cash benefit after recurring costs	$\$1,400 - \$950 = \$450$
Year-one gross cash benefit	$12 \times \$1,400 = \$16,800$
Year-one cash costs	$\$4,500 + (12 \times \$950) = \$15,900$
Year-one net cash benefit	$\$16,800 - \$15,900 = \$900$
Year-one simple cash ROI on total cash costs	$\$900 \div \$15,900 = 5.7\%$, rounded
Simple cash payback on initial outlay	$\$4,500 \div \$450 = 10$ months

The payback calculation assumes the initial outlay occurs at inception and the full monthly net benefit begins immediately and stays constant. If adoption ramps or benefits are delayed, use monthly cash flows instead. If recurring net cash benefit is zero or negative, this simple formula provides no positive payback period.

State what the return excludes.

This is not a discounted return, IRR, or accounting profit. It excludes taxes, financing, resale proceeds, and the separately disclosed 60 hours of setup opportunity cost. It is not a projection of actual bank results.

Avoided future hiring or spending is a forecast against a documented counterfactual, not a booked expense reduction. Keep it separate from realized savings and confirm that the canceled expense was credible, necessary, and not merely deferred or shifted elsewhere.

10 / TEST THE ASSUMPTIONS

A stronger decision survives a less tidy scenario.

Vary the drivers that could change the result. More decimal places cannot compensate for uncertain adoption, review burden, or task volume.

The table models three different operating conditions, not confidence intervals or predictions. Each uses 300 eligible monthly requests and a 30-minute manual baseline. Fallback percentage is always a share of attempted requests. Eligible requests without an AI attempt remain in the manual workload.

Monthly assumption or result	Downside	Base	Stronger
AI adoption among eligible requests	60%	80%	90%
Fallback among attempted requests	20%	10%	10%
Minutes per assisted completion	20	16	14
Total minutes per fallback request	40	38	38
Separate administration hours	8.0	6.0	5.0
Total proposed staff hours	140.0	108.8	93.8
Net recurring capacity hours	10.0	41.2	56.2
Useful hours at 75% redeployment	7.50	30.90	42.15
Useful capacity value at \$52/hour	\$390.00	\$1,606.80	\$2,191.80
Proxy less \$950 external monthly cost	-\$560.00	\$656.80	\$1,241.80

The last line mixes a noncash labor-value proxy with an external cash cost. It is a capacity-value comparison, not savings, profit, or cash ROI. It excludes initial cash investment and setup opportunity cost. Without a verified cash benefit, all three scenarios still add \$950 of recurring monthly cash spending.

Challenge one variable as well as the whole scenario

The three scenarios change several inputs together, so they do not identify which one matters most. Also vary adoption, fallback, reviewer minutes, useful redeployment, and usage cost individually while holding other inputs fixed. Test lower demand: overhead may remain even when there are fewer requests to support it.

At the assumed \$52 rate, covering \$950 of recurring external cost on a capacity-value basis requires 18.27 usefully redeployed hours per month, rounded. That threshold excludes setup and is not cash break-even. If the work has a different role mix or practical destination, the relevant value can change materially.

For a material investment, have Finance assess realistic monthly cash flows and, where appropriate, discounted value. Keep serious quality or control failures outside an average-benefit trade-off; a positive modeled return does not override an unacceptable use.

11 / MAKE THE DECISION

Write a memo that another leader can challenge.

A useful decision memo shows the boundaries, the evidence, and the open questions. It is more valuable than a single productivity percentage.

Decision-memo section	What the sponsor should be able to explain
Work and baseline	The complete task, eligible workload, comparison method, sample size, task mix, and what the evaluation excludes.
Quality and boundaries	Acceptance criteria, observed defects by severity, manual fallbacks, relevant information controls, and unresolved issues.
Capacity and service	Net hours by role, who can use them, the intended destination, and actual service outcomes rather than only theoretical availability.
Financial view	Initial and recurring cash costs, staff opportunity cost, realized versus forecast benefits, sensitivity, and no double counting.
Decision and ownership	Continue, revise, pause, or propose expansion; the approving owners, evidence still needed, and a date for review.

Use the evidence to choose the next action

- Continue narrowly when the accepted work is useful, the controls fit the scope, and the operating burden is justified by the chosen benefit. A service-led decision need not pretend to be a cash-savings decision.
- Revise and retest when a specific problem is fixable, such as unclear intake, an outdated source collection, or a review bottleneck. Record the change so the next comparison remains interpretable.
- Pause or stop when unacceptable information handling, misleading outputs, unmanageable review effort, or unresolved commitments remain. Keep the manual route available.
- Propose expansion only after describing the additional workload, users, information, and permitted actions. Evidence for one scope is not approval for another.

Set an operating review that checks demand, adoption, fallback, quality, total staff effort, usage charges, and the actual destination of capacity. Retest after material changes to the service, sources, workflow, permissions, or staff responsibilities. The community-bank guide treats ongoing monitoring and termination planning as part of the relationship, not just the purchase. [2]

The strongest claim may be a narrow one.

A statement such as “The evaluated process reduced handling effort for this task under these conditions; cash savings are not yet demonstrated” is more credible than “AI saves the bank money” when the broader evidence does not exist.

12 / SUPPORTING WORKSHEET

Put the method to work on one assignment.

Complete this page after reading the guide. Use one common period and retain an evidence reference or owner for every material assumption.

Define the assignment	Your entry
Workflow and accepted deliverable	_____
Business owner / reviewer / Finance owner	_____
Period, task categories, and quality criteria	_____
Approved information and excluded actions	_____

Record the common-period inputs	Value	Evidence / owner
N: All eligible requests	_____	_____
A: Requests on which AI was attempted	_____	_____
S: Assisted completions / F: manual fallbacks	____ / ____	_____
b: Baseline staff minutes per request	_____	_____
s: Assisted-completion staff minutes	_____	_____
f: Fallback minutes, including attempt + manual	_____	_____
H: Separate recurring admin hours	_____	_____
Useful redeployment % / rate by role	____ / ____	_____
Initial cash / recurring cash per period	____ / ____	_____
Initial staff hours / verified cash benefits	____ / ____	_____

Check: $A = S + F$; $N - A$ stays manual.

Net hours = $(N \times b) \div 60 - [S \times s + F \times f + (N - A) \times b] \div 60 - H$.

This compact formula assumes the same baseline rate across requests and that every attempted request is either an assisted completion or a manual fallback. For different task categories, compute each separately and sum; deduct shared administration only once. Add separate paths for canceled, unresolved, or other outcomes rather than forcing them into a false completion count.

Before stating the result

- Comparable quality and workload are supported; failed attempts and downstream corrections are included.
- Capacity, service, cash, and opportunity cost remain distinct; no hour or expense is counted twice.
- Observed results, estimates, and targets are labeled; serious exceptions remain visible.
- Decision, owner, evidence gaps, and next review are recorded: _____.

This is a supporting calculation aid, not a compliance approval or a substitute for Finance review. The Word edition is editable; the PDF worksheet is intended for printing or annotation, not automatic calculation.

NEXT STEP / SOURCES

Bring a workflow, not a promised return.

Start with the work staff need to complete and the evidence that would make the decision easier.

A useful discussion with Providensi begins with a general workflow description, the intended deliverable, approximate volume, current handling effort, relevant information requirements, and the people who will review the result. You do not need to supply customer records or credentials to explain the business problem.

Discuss one measurable banking workflow with Providensi.

Choose whether the primary goal is useful staff capacity, better service, or a documented cash outcome. Bring synthetic or appropriately approved examples and the questions your business, Finance, technology, and risk teams need answered. Scope, delivery, pricing, and commitments require separate agreement.

Reading notes and authoritative references

References checked October 6, 2026. Bracketed citations support the general evaluation and third-party oversight principles described in the text. The measurement protocol, calculations, fictional example, scenario table, decision memo, and worksheet are Providensi's educational synthesis. They are not agency-issued formulas, regulatory thresholds, benchmarks, or observed customer results.

[1] Artificial Intelligence Risk Management Framework (AI RMF 1.0)

NIST • January 26, 2023

A voluntary framework. The MAP, MEASURE, and MANAGE functions address intended context, anticipated benefits and costs, performance evaluation, uncertainty and limitations, deployment decisions, and monitoring. See MAP 1.4 and 3.1–3.2; MEASURE 1.1, 2.3–2.5, and 4.1; and MANAGE 1.1 and 2.1.

[2] Third-Party Risk Management: A Guide for Community Banks

Federal Reserve, FDIC, and OCC • May 2024; issued May 3, 2024

Printed pages 1–2 explain continuing bank responsibility and the guide's nonbinding status. Pages 5–6 discuss strategic fit, benefits, direct and indirect costs, and planning resources. Pages 14–18 address ongoing monitoring and termination. The guide is not a checklist, substitute for the underlying guidance, or regulatory safe harbor.

Scope and limitations

This guide is educational, not legal, accounting, investment, or regulatory advice. It does not establish regulatory approval, security assurance, or an expected financial return. A bank's responsible teams determine applicable obligations, acceptable use, accounting treatment, approval requirements, and the evidence needed for the actual arrangement.

Companion in this series: A Bank Leader's Guide to Practical AI. That guide introduces supervised workflows, information boundaries, and pilot readiness; this guide develops the value-measurement method in more detail. Page references follow the designed PDF; editable Word pagination may vary by application.