

A Bank Leader's Guide to Practical AI

How to choose useful work, establish clear boundaries, and
evaluate the results

Less preparation. More capacity for customers.

AI becomes useful when it helps people complete work they can trust and review. This guide explains how to move from an interesting demonstration to a defined banking assignment, with accountable people, appropriate information handling, and evidence of value.

01

Recognize a sensible first
workflow

02

Distinguish assistance from
authority

03

Evaluate quality, effort, and
operating fit

For leaders of small and midsized banks

An educational guide with practical examples and a short readiness checklist. No technical background required.

01 / THE LEADERSHIP DECISION

Start with work, not with a model.

A useful first question is not “Which AI should we buy?” It is “Which recurring task would become more valuable if our people spent less effort preparing it?”

Consider the effort behind a policy question. Someone locates the document, checks the version, finds the relevant passage, drafts an answer, and asks a colleague to review it. AI may help with parts of that preparation. The policy owner still determines what is authoritative, and the bank remains responsible for how the answer is used.

This shifts the evaluation from a broad promise of intelligence to an observable result: a draft answer with the right sources, a usable evidence index, or an exceptions list that helps a reviewer find the problem. The benefit is not simply more output. It is less avoidable effort between a request and a sound decision.

Three questions for the sponsor

- **What work should improve?** Name the people, recurring task, current difficulty, and useful deliverable. “Explore AI” is an interest, not an assignment.
- **What must remain under control?** Identify information, system access, external actions, and decisions that the assignment must not take over.
- **What would count as improvement?** Compare quality, total staff effort, rework, operating cost, and service outcomes against a baseline.

A demonstration is a starting point, not proof of readiness.

A polished answer does not establish dependable source use, appropriate access, production capacity, or acceptable contractual terms. Those are separate questions, and each needs its own evidence.

Use this guide as a shared conversation

Business leaders can use it to define value. Technology and security teams can use it to clarify the environment. Risk, compliance, and vendor-management colleagues can use it to identify the reviews that apply. The interagency community-bank guide emphasizes that engaging a third party does not remove the bank’s responsibilities. [1]

Understand the work: pages 3–5

Establish the operating boundaries: pages 6–8

Evaluate and decide: pages 9–11

Providensi and reference notes: page 12

Scope: internal, supervised preparation work. This guide is educational, not legal advice, a complete control framework, a compliance certification, or a substitute for your bank’s policies. All worked examples are fictional.

02 / A CLEARER MENTAL MODEL

Answers, work, and actions are different.

Two tools may look like the same chat window while having very different access and consequences.

A generative AI system produces new content, such as a summary, explanation, or draft. A connected assistant can also use tools to retrieve files, create documents, or interact with another system. When those capabilities are combined, the important question becomes what the complete workflow is allowed to do, not merely how well it writes.

Level of work	Banking example	What to establish
Explain	Summarize an approved procedure for a staff member.	Which sources are authoritative, whether the staff member may access them, and who verifies the interpretation.
Prepare	Build a draft comparison or evidence index from supplied files.	Where files are processed, whether originals stay unchanged, and how findings trace back to source material.
Act	Update a record, send a message, or submit a document.	Explicit authority, tested permissions, recipient checks, review points, and recovery or correction procedures.

A source-backed answer is still an answer to check

Some systems retrieve relevant material from an approved collection before generating a response. This is often called retrieval-augmented generation. It can help ground a response in supplied information, but retrieval can miss a relevant passage, select the wrong version, or return material the requester should not see. A citation can also be inaccurate or fail to support the conclusion. [3]

Ask the reviewer to open the cited passage and compare it with the answer. Check the document’s approval status and scope, not just its date. A later draft is not automatically the effective policy. An accurate quotation can still be applied to the wrong situation.

Read-only does not mean disclosure-free.
An assignment that cannot change a source system may still process information elsewhere or produce a file that someone can share. Evaluate reading, processing, saving, and sending as distinct permissions.

For a first pilot, separate useful preparation from consequential action. Keep credit decisions, account handling, payments, customer communications, and regulatory submissions outside the initial scope unless a separately approved use case and its required controls have been established. This is a conservative starting approach, not a claim that a read-only workflow needs no review.

03 / SELECTING THE FIRST ASSIGNMENT

Choose work with a checkable result.

Start where the inputs can be controlled, a knowledgeable person can verify the output, and a mistake can be contained before anyone relies on it.

Policy and procedure support

Give the system a small set of approved documents and a defined staff question. Ask for a draft answer, source passages, version information, and unresolved conflicts. A policy owner checks whether the interpretation is valid for the situation. The useful result is less searching and assembly, not a new authority to decide bank policy.

Examination-evidence preparation

Use a synthetic request list and approved sample materials to prepare an evidence index. Each request should link to an actual source, its owner, and its status. Missing approval records stay missing. The bank team decides whether evidence is sufficient and whether anything may be submitted. An organized packet is not an examination outcome.

Document comparison and reconciliation

Compare synthetic reports using explicitly agreed reference fields. Ask for source locations, unmatched entries, missing values, and differences that need review. Recalculate important totals with a deterministic method, such as a spreadsheet formula, and verify extracted data against the originals. A plausible match is not a verified match.

Choose the narrowest useful version

A procedure comparison is easier to evaluate when it covers one topic and a known source set than when it searches every shared drive. A document check is easier to review when the required fields are named in advance. Narrow scope reduces the number of things you must understand at once; it does not make the work automatically safe.

Avoid choosing solely on task volume. A frequent task may still involve restricted information, ambiguous judgment, or difficult system dependencies. Prefer an assignment for which a domain expert can explain what a good answer contains, what a wrong answer looks like, and when the system should stop and ask.

Try this selection test.

Can we supply appropriate examples, define the expected result, identify the reviewer, and keep mistakes away from live decisions? If not, simplify the assignment before expanding the technology.

These are candidate workflows to evaluate, not reported customer results or promises of compatibility with a particular bank system.

04 / A WORKED POLICY-SUPPORT EXAMPLE

A useful answer can leave a question open.

The following scenario is entirely fictional. The procedure wording, documents, and sample response are invented teaching materials, not bank guidance or observed product output.

A staff member asks who may approve a procedure exception. The approved evaluation folder contains Procedure A, marked effective, and a later Procedure B, visibly marked draft. No record approving Procedure B has been supplied. The task is to prepare a source-based comparison for a policy owner, not to approve the exception.

Synthetic source	What the supplied material says	What it does not prove
Procedure A, section 4	An effective procedure names the branch manager as the exception reviewer.	That every type of exception falls within the section's scope.
Procedure B, section 2	A later draft proposes review by an operations manager.	That the draft was approved or replaced Procedure A.
Source register	Both files and their locations are listed; approval evidence for B is absent.	That an approval exists elsewhere or may be inferred.

Example assignment

Compare the supplied procedures for this question. Cite the section and relevant wording, distinguish effective material from drafts, and list what the sources do not establish. Do not approve an exception, search other accounts, contact anyone, or change the source files.

What a useful draft would contain

A useful draft would identify the reviewer named in Procedure A and separately describe the proposal in Procedure B. It would explain that naming a reviewer does not itself establish who may approve an exception. The policy owner must confirm approval authority, the applicable exception type, and whether an approved revision exists. The draft would not quietly choose the newer file because its date looks more recent.

The human reviewer then opens both sections, checks their status and applicability, resolves the outstanding question, and records the accepted answer. If the answer changes, the reviewer keeps the correction with the source reference. That record can improve the next assignment and reveal recurring weaknesses.


What this teaches: source retrieval, document status, interpretation, and approval are separate steps. A confident answer that erases those distinctions can be less useful than a careful draft that exposes the missing evidence. NIST identifies confident false content and misleading citations as generative-AI risks. [3]

05 / PEOPLE AND RESPONSIBILITIES

Design the handoff, not just the prompt.

An internal AI coordinator can make access and review more deliberate. The coordinator is an operating role, not a substitute for technical controls or bank accountability.

One approach is to start with a trained coordinator who receives staff requests, checks their suitability, prepares the assignment, and returns the result to a designated reviewer. This can help a bank learn before considering wider access. It should be evaluated against the bank’s staffing, service needs, and existing governance rather than adopted as a universal model.



Responsibility	The decision that stays with people
Business sponsor	Defines the problem, intended benefit, boundaries, and conditions for expansion.
Coordinator or operator	Uses approved sources and accounts, follows the scope, and escalates exceptions rather than improvising new authority.
Domain reviewer	Checks evidence, calculations, applicability, and completeness before accepting the deliverable.
Technology and security	Confirm permissions, processing paths, logging, restrictions, and interruption or incident handling.
Risk, compliance, and vendor management	Determine applicable review expectations, assess evidence and commitments, and challenge unresolved assumptions.

These are responsibilities, not a required staffing chart. One person may hold more than one role where appropriate, but a bank should address independence, expertise, and conflicts when deciding who reviews consequential work. Name a backup reviewer and a route for urgent questions so the process does not depend on one person’s availability.

Make review real

“Human in the loop” is incomplete unless someone has the time, knowledge, evidence, and authority to disagree. Give reviewers the source material, show unresolved items separately, and allow them to reject the result. Do not measure success in a way that rewards quick approval over careful challenge. NIST describes over-reliance on automated output as a risk in human-AI interaction. [3]

Watch for a hidden bottleneck.


Compare drafting time saved with added coordinator and reviewer effort. If that extra effort uses up or exceeds the saving, there is no net capacity gain. Measure the whole handoff, including corrections and escalations.

06 / INFORMATION BOUNDARIES

“Private” needs an information map.

Where a file is stored, where it is processed, who can retrieve it, and where the answer can be sent are different questions.

A desktop application can keep a project file on the bank’s computer while sending relevant content to an external service for an operation. A dedicated deployment can use bank-controlled infrastructure and still depend on online research, updates, support access, or optional connections. A local label alone does not settle those paths.



At each step, identify the location, permitted users, service providers, possible transfers, and applicable retention rules. Include logs and backups, not just the original file.

Question	Evidence that helps answer it
What enters the assignment?	A source inventory and data classification covering attachments, prompts, copied text, credentials, and retrieved passages.
Where is it processed?	A deployment-specific information-flow diagram and an inventory of the services used for each operation.
Who can access the material?	Role and permission definitions, including source access, administration, and temporary support access, backed by tests.
What is retained or transferred?	Applicable terms and tested settings for outputs, logs, indexes, backups, exports, deletion, and external sharing.

Match the environment to the workload

Bank-hosted, private-cloud, and hybrid approaches create different operating arrangements. A hybrid approach may use local work for one task and an external service for another. Do not permit a sensitive document to follow the second path merely because the first path was approved. Evaluate the actual configuration, responsibilities, costs, and restrictions.

Start with synthetic or appropriately approved examples. Removing a name does not necessarily remove sensitivity or the possibility of identifying someone. Access should be enforced by the relevant systems, not solely by asking the assistant to behave. A restriction must still hold when the request is phrased differently or an unexpected file appears.

One boundary worth making explicit

Permission to read a document is not permission to upload it, email it, or include its contents in an outside search. Agree on each use and test the controls that enforce it. [1, 3]

07 / VENDOR AND CONTROL EVIDENCE

Ask about the service you will actually use.

Due diligence becomes more useful when the evidence request describes a workload, information set, access level, and operating environment.

An assistant that summarizes approved public material is not the same engagement as a system that reads customer records or updates a production application. Describe the proposed use before comparing vendors. The interagency community-bank guide recommends oversight proportionate to the relationship’s risk and covers the full relationship life cycle. [1]

Stage	Practical question for this AI service
Planning	What business activity, information, and system dependencies are involved, and what is deliberately excluded?
Due diligence and selection	Which evidence covers this service, its subcontractors, security, operating capability, and limitations?
Contract negotiation	Which data-use, support, incident, change, and exit commitments are actually agreed?
Ongoing monitoring	How will the bank learn about changes, check performance, and follow up on exceptions?
Termination	Can the bank retain usable work, revoke access, address retained data, and continue the process elsewhere?

Separate the claim from the evidence

For each important claim, ask what was examined, in which environment, when, by whom, and with what exceptions. A report about one service may not cover the proposed deployment. A test of ordinary inputs may not address unauthorized access. A catalog entry for an application connection does not prove that a particular core-banking workflow is supported.

Security review should consider ordinary controls and AI-specific failure modes. One example is prompt injection: instructions hidden in material the system encounters can try to redirect the assignment. Test how supplied content is treated, restrict available actions, and inspect attempted or completed effects rather than relying only on a reassuring response. NIST describes this risk in its generative-AI profile. [3]

Missing evidence is an open item, not an invented approval.

Record the requirement, available evidence, gap, owner, and decision needed. A checklist can organize that record; checking every box does not establish regulatory approval or create a safe harbor. The agencies’ community-bank guide makes a similar limitation explicit. [1]

08 / A CONTROLLED PILOT

Test usefulness and boundaries together.

Define what the pilot may do, what a good result contains, and what would stop the evaluation before the first test runs.

- 1. Agree on one assignment.** Name the requester, approved input set, expected deliverable, reviewer, and excluded actions. State whether the system may only read and prepare, and which tools or accounts are unavailable.
- 2. Establish the reference answer.** Have a qualified reviewer identify the expected findings and acceptable treatment of uncertainty. Where judgment is involved, document the reasoning and discuss disagreements rather than treating an existing answer as infallible.
- 3. Use a representative test set.** Include typical work, difficult examples, and boundary cases. Keep some examples separate from those used to refine instructions so the pilot does not merely reproduce rehearsed answers.
- 4. Inspect results and side effects.** Check source accuracy, completeness, corrections, and human effort. Also inspect files, access attempts, messages, and system changes where relevant. “Done” in a chat is not proof that the intended action completed.
- 5. Decide before expanding.** Review the evidence with the responsible owners. Record acceptance, remediation, or rejection, and define what must be retested after a material change.

Include this test	Look for this behavior
Missing or conflicting source material	Identifies the gap and cites the conflict rather than inventing a resolution.
A requester without access to a source	Cannot retrieve or disclose that source through an alternative request.
An instruction embedded in a document	Does not treat source text as authority to change the assignment or send information.
An interruption or stop request	Shows what completed, what did not, and whether resuming could repeat an action.

Set acceptance thresholds according to the use case and consequences. Record serious failures separately instead of hiding them in an average score. Limited successful tests support only the conditions tested; they do not prove universal accuracy or security. [2, 3]

This is a proposed evaluation method, not an official examination checklist. Your bank determines the required approvals and evidence for the intended use.

09 / THE BUSINESS CASE

Measure capacity before claiming savings.

A faster draft is useful only if it improves the complete job. Count preparation, review, corrections, support, and operating cost.

Compare the same type of assignment under the existing and proposed processes. Keep volume, complexity, and quality expectations comparable. Record how time was measured, who performed the work, and what the sample excludes. A stopwatch on the AI response alone misses the people who prepared inputs or repaired the result.

A practical capacity calculation
Net recurring staff capacity = comparable baseline staff hours – AI-assisted preparation, supervision, review, and correction hours – recurring administration hours. Record one-time setup and training effort separately. Use a consistent period and do not subtract the same activity twice.

Measure	How to interpret it
Staff effort	Count time across every person involved. Separate active staff time from elapsed turnaround and unattended machine time.
Quality and rework	Record missing evidence, incorrect citations, extraction errors, reviewer corrections, and rejected outputs. Keep severe errors visible.
Operating cost	Include relevant service usage, licensing, infrastructure, support, monitoring, and administration. Explain one-time versus recurring costs.
Customer or business capacity	Record whether staff used reclaimed time for customer service, follow-up, or other useful work, rather than assuming it happened.

A worked interpretation, without invented ROI

Imagine an evidence-index pilot that reduces file assembly but creates more checking for the reviewer. First combine both people’s effort. Then include the recurring work needed to keep source files and instructions current. If that total is lower at comparable quality, the pilot may have reclaimed capacity. If the total is unchanged, the faster draft has not yet improved labor efficiency.

Reallocated time is not automatically lower payroll or a better efficiency ratio. A cash-savings claim needs an actual expense reduction or a documented cost avoided, with a credible baseline. A revenue claim needs observed business outcomes and a defensible connection to the workflow. Do not value the same reclaimed hour twice as both a labor saving and a separate benefit without explaining the distinction.

No customer savings or productivity benchmarks are asserted here. This is a measurement approach. Choose the evidence and decision criteria before the pilot so a compelling anecdote does not become the whole business case.

10 / THE EXPANSION DECISION

Expand the evidence before the authority.

A successful narrow pilot is a reason to consider the next scope, not permission to add every source, user, or connected action.

Adding a new document collection, recipient, integration, or user group can change the risk even when the software stays the same. Moving from draft preparation to sending or record updates is a particularly important change. Revisit the information map, permissions, review responsibilities, and test set before making that move.

Decision	What would support it
Continue narrowly	The agreed result is useful and reviewable, boundaries hold in the tested conditions, and the operating burden is acceptable.
Revise and retest	A specific source, instruction, interface, or process problem can be addressed without concealing unresolved failures.
Pause or stop	Unacceptable disclosure, unauthorized action, unreliable evidence, unmanageable review burden, or unresolved operating commitments remain.
Propose expansion	The sponsor identifies an additional benefit, the relevant owners approve the new scope, and evidence supports the changed conditions.

Keep a practical operating record

Retain the approved scope, source versions, configuration or release, test results, limitations, owners, and decisions in the bank’s chosen record. Set who will review exceptions, vendor changes, source changes, and recovery exercises. Agree on triggers for retesting and a usable manual alternative. Ongoing monitoring and termination planning belong to the relationship, not only its purchase decision. [1]

A short readiness check

- A specific internal outcome and a business owner are named.
- Inputs, processing paths, retention, and excluded information are understood.
- Permitted reads, changes, and sends are separately defined and tested.
- A qualified reviewer has time, evidence, and authority to challenge results.
- Representative quality and boundary tests have agreed acceptance criteria.
- Total staff effort and operating costs have a comparable baseline.
- Support, change, interruption, incident, and exit responsibilities are addressed.
- The bank has recorded what may proceed, what remains open, and what triggers a stop.

This readiness check is a discussion aid. It does not grant access, authorize production use, establish compliance, or replace required bank review. An unresolved item should remain visible with an owner and next action.

NEXT STEP / SOURCES

Bring one useful workflow to the table.

Start the conversation with the work your staff need to complete, not with a request to replace an entire banking system.

Providensi offers separately scoped bank-hosted LLM deployments and an evaluation approach centered on approved knowledge, customer-data requirements, and work bank staff can review. General Thaddeus capabilities include research, document preparation, supported browser-workflow reuse, and project continuity. Inclusion, access, local execution, hardware fit, and service commitments must be confirmed for the bank's agreed deployment.

Discuss your bank's first AI assignment with Providensi.

Bring a general workflow description, the deliverable you need, the people who will review it, and your information requirements. Begin with synthetic or approved examples, not customer records or credentials. The next step is a scoped evaluation discussion, not a promise of deployment or results.

Reading notes and authoritative references

References checked October 6, 2026. Source-backed principles are identified by bracketed numbers. The example assignments, decision tables, and readiness check are Providensi's educational synthesis, not agency-issued requirements or observed customer results.

[1] Third-Party Risk Management: A Guide for Community Banks

Federal Reserve, FDIC, and OCC • May 2024

See printed pages 1–5, 8–21. Explains the relationship life cycle and proportionate oversight; it is not a regulatory safe harbor or a replacement for the underlying guidance.

[2] AI Risk Management Framework (AI RMF 1.0)

NIST • January 26, 2023

A voluntary framework organized around Govern, Map, Measure, and Manage. Using it is not regulatory certification.

[3] Artificial Intelligence Risk Management Framework: Generative Artificial Intelligence Profile

NIST AI 600-1 • July 26, 2024

See sections 2.2, 2.4, 2.7–2.9 and 2.12 for false content, privacy, human reliance, information integrity, security, and component risks.

[4] Supervisory Guidance on Model Risk Management

Federal Reserve, FDIC, and OCC • April 17, 2026

Page 3, footnote 3 excludes generative and agentic AI from this guidance's scope and points to broader bank risk-management and governance practices for tools outside it. That exclusion does not remove other applicable obligations.

Important scope distinction: do not describe every generative-AI pilot as governed by the agencies' 2026 model-risk document. Ask the responsible bank teams which policies and obligations apply to the actual use. [4] This guide is not a legal opinion, security assessment, certification, or guarantee of performance.